

# HPC & Big Compute with Microsoft Azure

Karl Podesta  
Global Black Belt Team – HPC & Big Compute specialist  
[kapodest@microsoft.com](mailto:kapodest@microsoft.com)  
<http://microsoft.com/hpc>

# Mission

Our mission is to enable researchers, engineers, analysts, designers, developers, and data scientists to achieve **radically better results** and **faster answers** to complex problems, by making it easy to do simulation and parallel computing **in the cloud, at hyper-scale**.

Azure empowers these technical experts to **think at 10x or 100x the current scale**, to work faster, better, and in new ways that before the cloud **we could only dream about**.







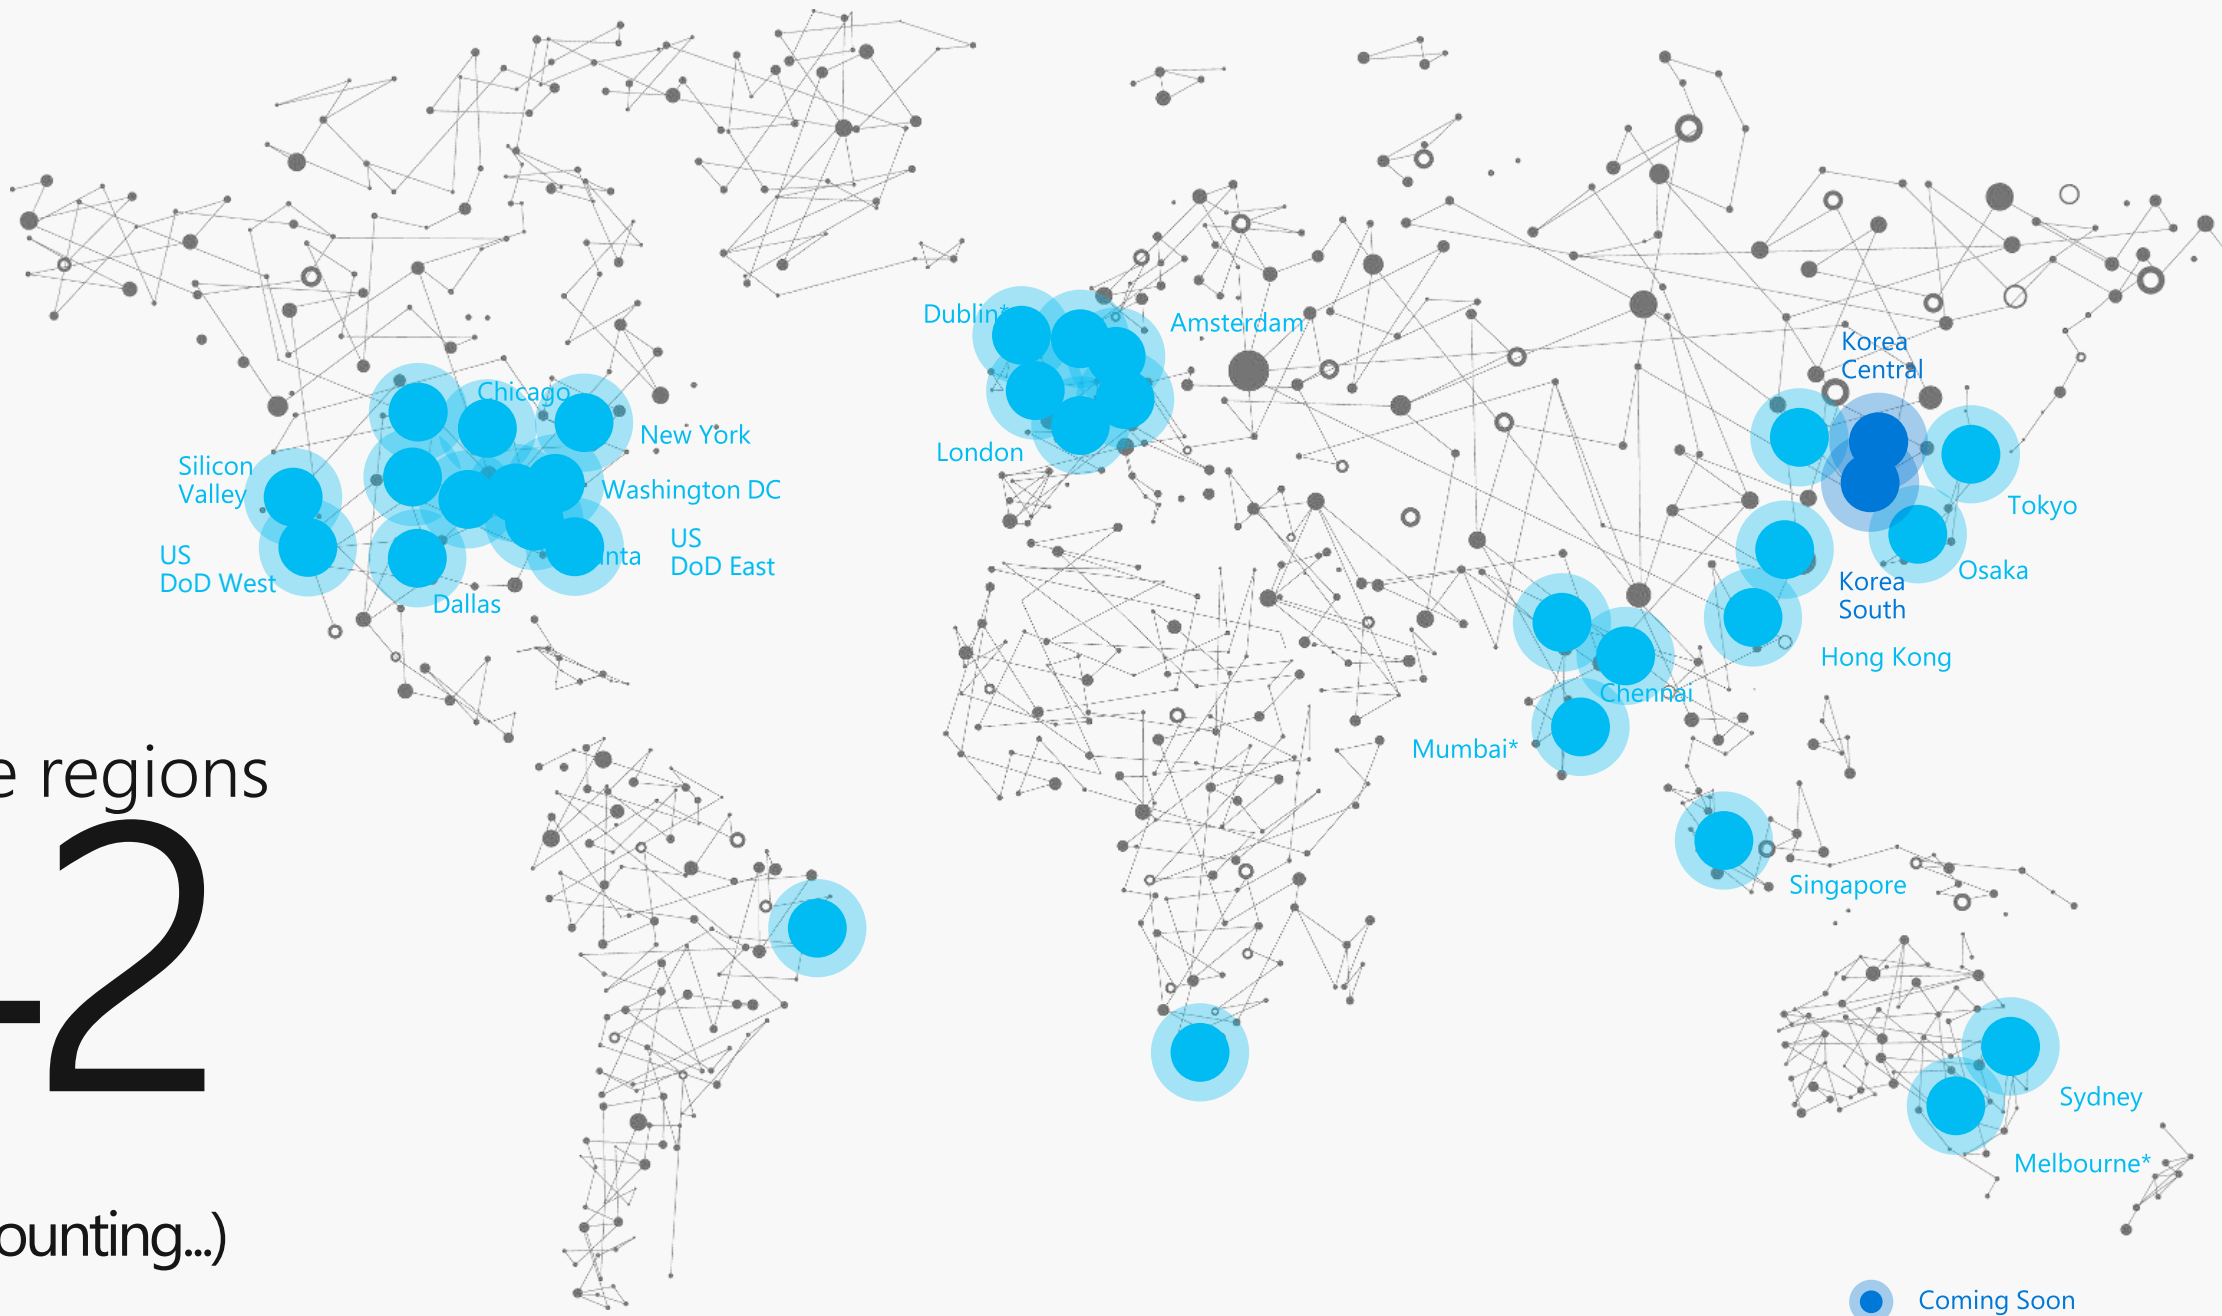
# Introduction to Azure



Azure regions







42

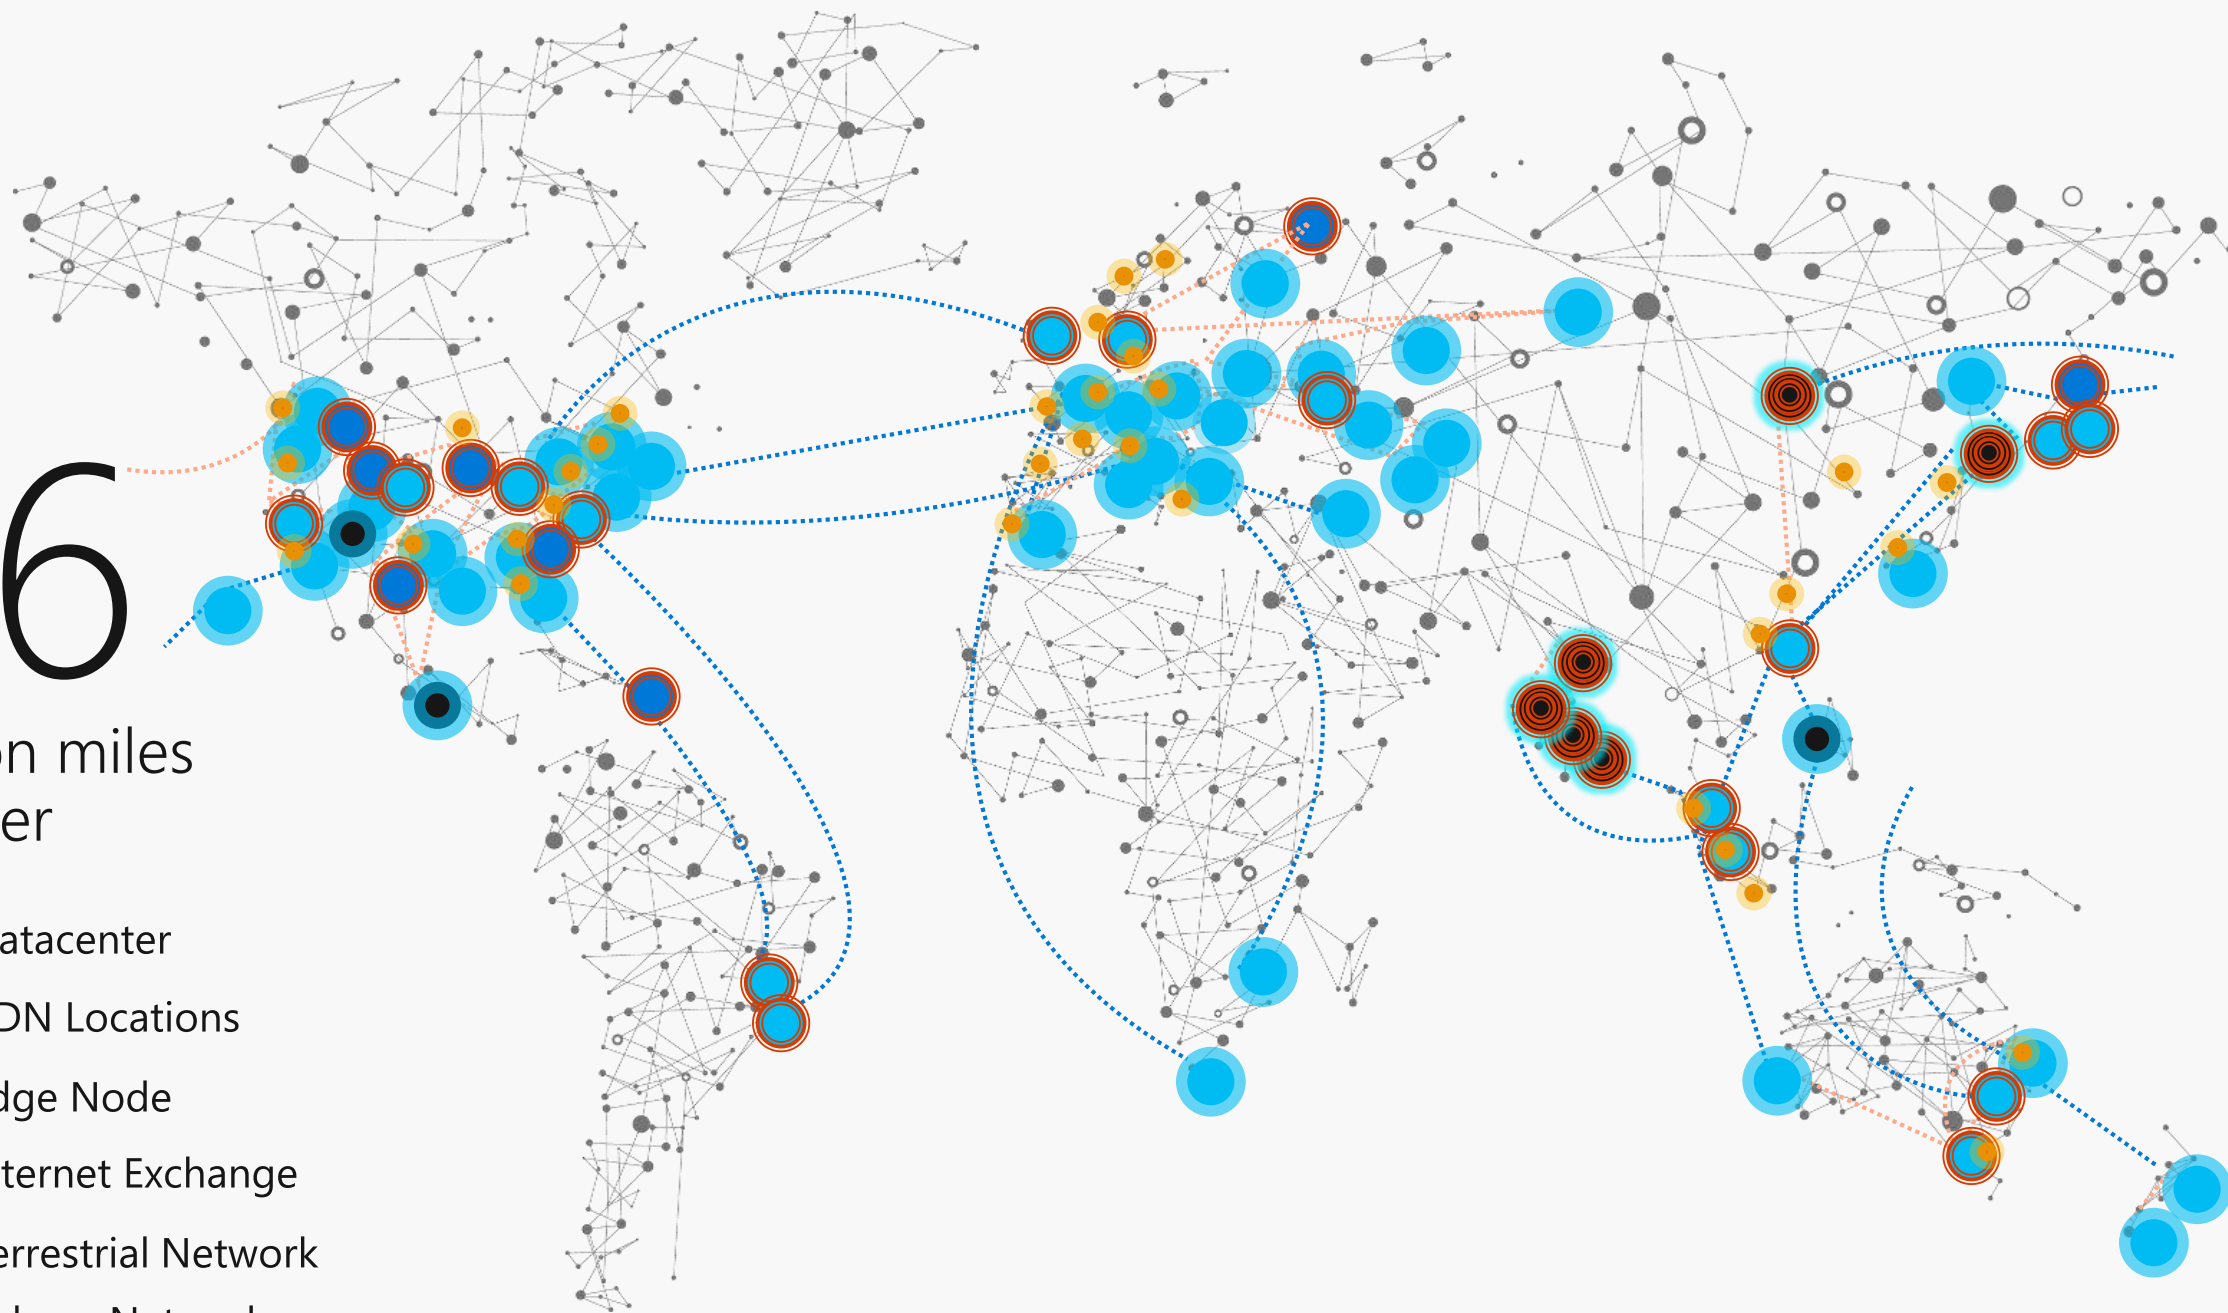
(and counting...)



# 1.6

million miles  
of fiber

-  Datacenter
-  CDN Locations
-  Edge Node
-  Internet Exchange
-  Terrestrial Network
-  Subsea Network



# ExpressRoute

Private connection to Azure

High-throughput

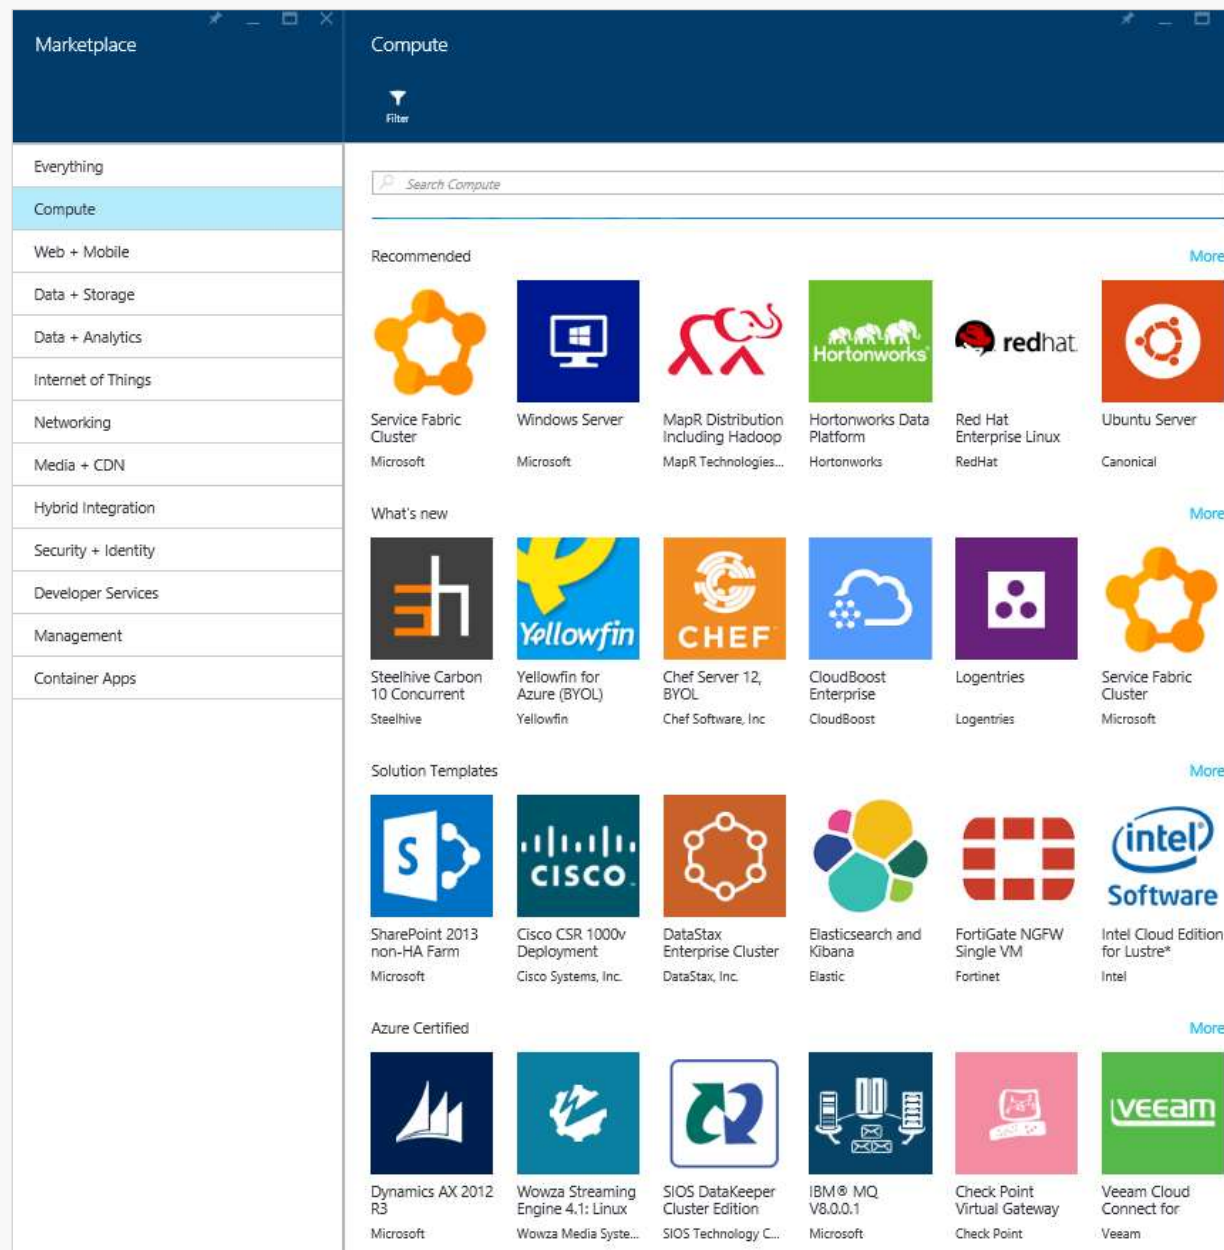
Predictable performance





# >4000

## Certified, pre- configured solutions in Azure Marketplace



# Azure covers 71 compliance offerings

Azure has the deepest and most comprehensive compliance coverage in the industry

## Global



ISO 27001



ISO 27018



ISO 27017



ISO 22301



SOC 1  
Type 2



SOC 2  
Type 2



SOC 3



CSA STAR  
Self-Assessment



CSA STAR  
Certification



CSA STAR  
Attestation

## US Gov



Moderate  
JAB P-ATO



High  
JAB P-ATO



DoD DISA  
SRG Level 2



DoD DISA  
SRG Level 4



DoD DISA  
SRG Level 5



SP 800-171



FIPS 140-2



Section  
508 VPAT



ITAR



CJIS



IRS 1075

## Industry



PCI DSS  
Level 1



CDSA



MPAA



FACT  
UK



Shared  
Assessments



FISC  
Japan



HIPAA /  
HITECH Act



HITRUST



GxP  
21 CFR Part 11



MARS-E



IG Toolkit  
UK



FERPA



GLBA



FFIEC

## Regional



Argentina  
PDPA



EU  
Model Clauses  
G-Cloud



UK



China  
DJCP



China  
GB 18030



China  
TRUCS



Singapore  
MTCS



Australia  
IRAP/CCSL



New  
Zealand  
GCIO



Japan My  
Number Act



ENISA  
IAF



Japan CS  
Mark Gold



Spain  
ENS



Spain  
DPA



India  
MeitY



Canada  
Privacy  
Laws




Privacy  
Shield



Germany IT  
Grundschutz  
workbook




Open and  
integrated




## Industry specific


### Life sciences



### Finance



### Manufacturing



## Software and hardware

### Languages



### Operating systems



### Infrastructure



# We Research and we Add Value

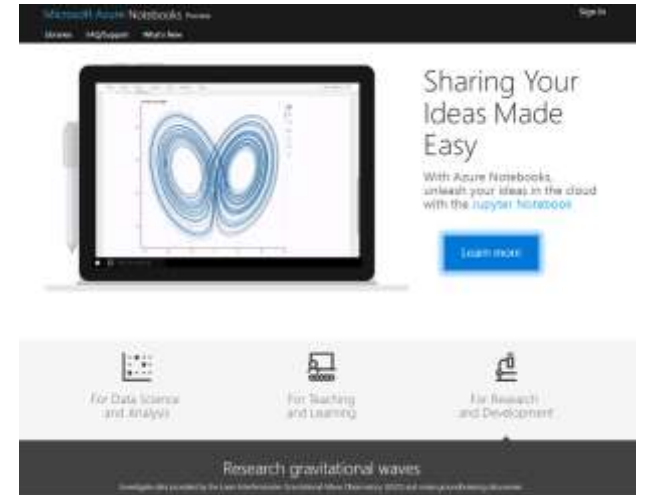
## Grants for non-profit research



# Microsoft Genomics Service



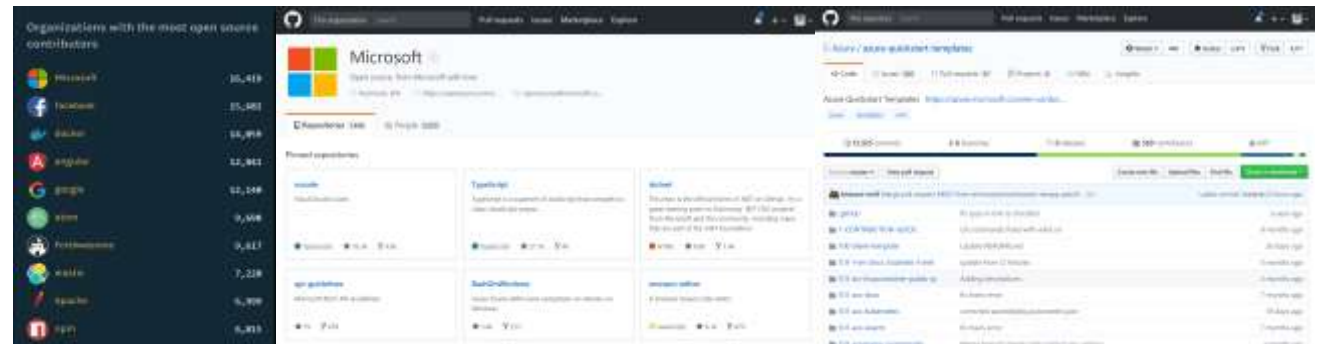
## Azure Notebooks



We will bring Quantum to everyone



## Top contributors on GitHub; Solutions/Recipes shared







# What is High Performance Computing (HPC)?





# Heard about Big Data? ... What about Big Compute?




How can our banks understand risk? Can we have safer cars?  
How will global weather change affect us? Can we find & cure disease?

# HPC is about doing even more...




Brady Holt ([https://commons.wikimedia.org/wiki/File:1997\\_Pontiac\\_Trans\\_Sport\\_SE\\_IIHS.jpg](https://commons.wikimedia.org/wiki/File:1997_Pontiac_Trans_Sport_SE_IIHS.jpg)), „1997 Pontiac Trans Sport SE IIHS“, <https://creativecommons.org/licenses/by/3.0/legalcode>



With HPC you can run more tests, analyze more scenarios, tackle more complex problems, explore new solutions, get faster answers...







# The new HPC Opportunity

Expand your HPC environment to the Cloud





# Elasticity




# What would you do with 100x the scale?

## Do more with hyper-scale:


- Service more users
- Run more projects
- Get results faster
- Run larger simulations
- Explore new insights (e.g., "What if?")

## Remove current limitations:

- Modify more parameters
- Analyze more complex models
- Visualize larger results
- Run more iterations
- Generate higher fidelity results
- Simulate longer periods of time



# Cloud value for Big Compute (HPC)



Scale



Elasticity



Pay for use



Reach & locality



Hardware selection




No infrastructure






Known costs



# Why InfiniBand RDMA matters?




# STAR-CCM+ Scalability to 1024 cores




<https://azure.microsoft.com/en-us/blog/availability-of-star-ccm-on-microsoft-azure/>

# Fluent Benchmark






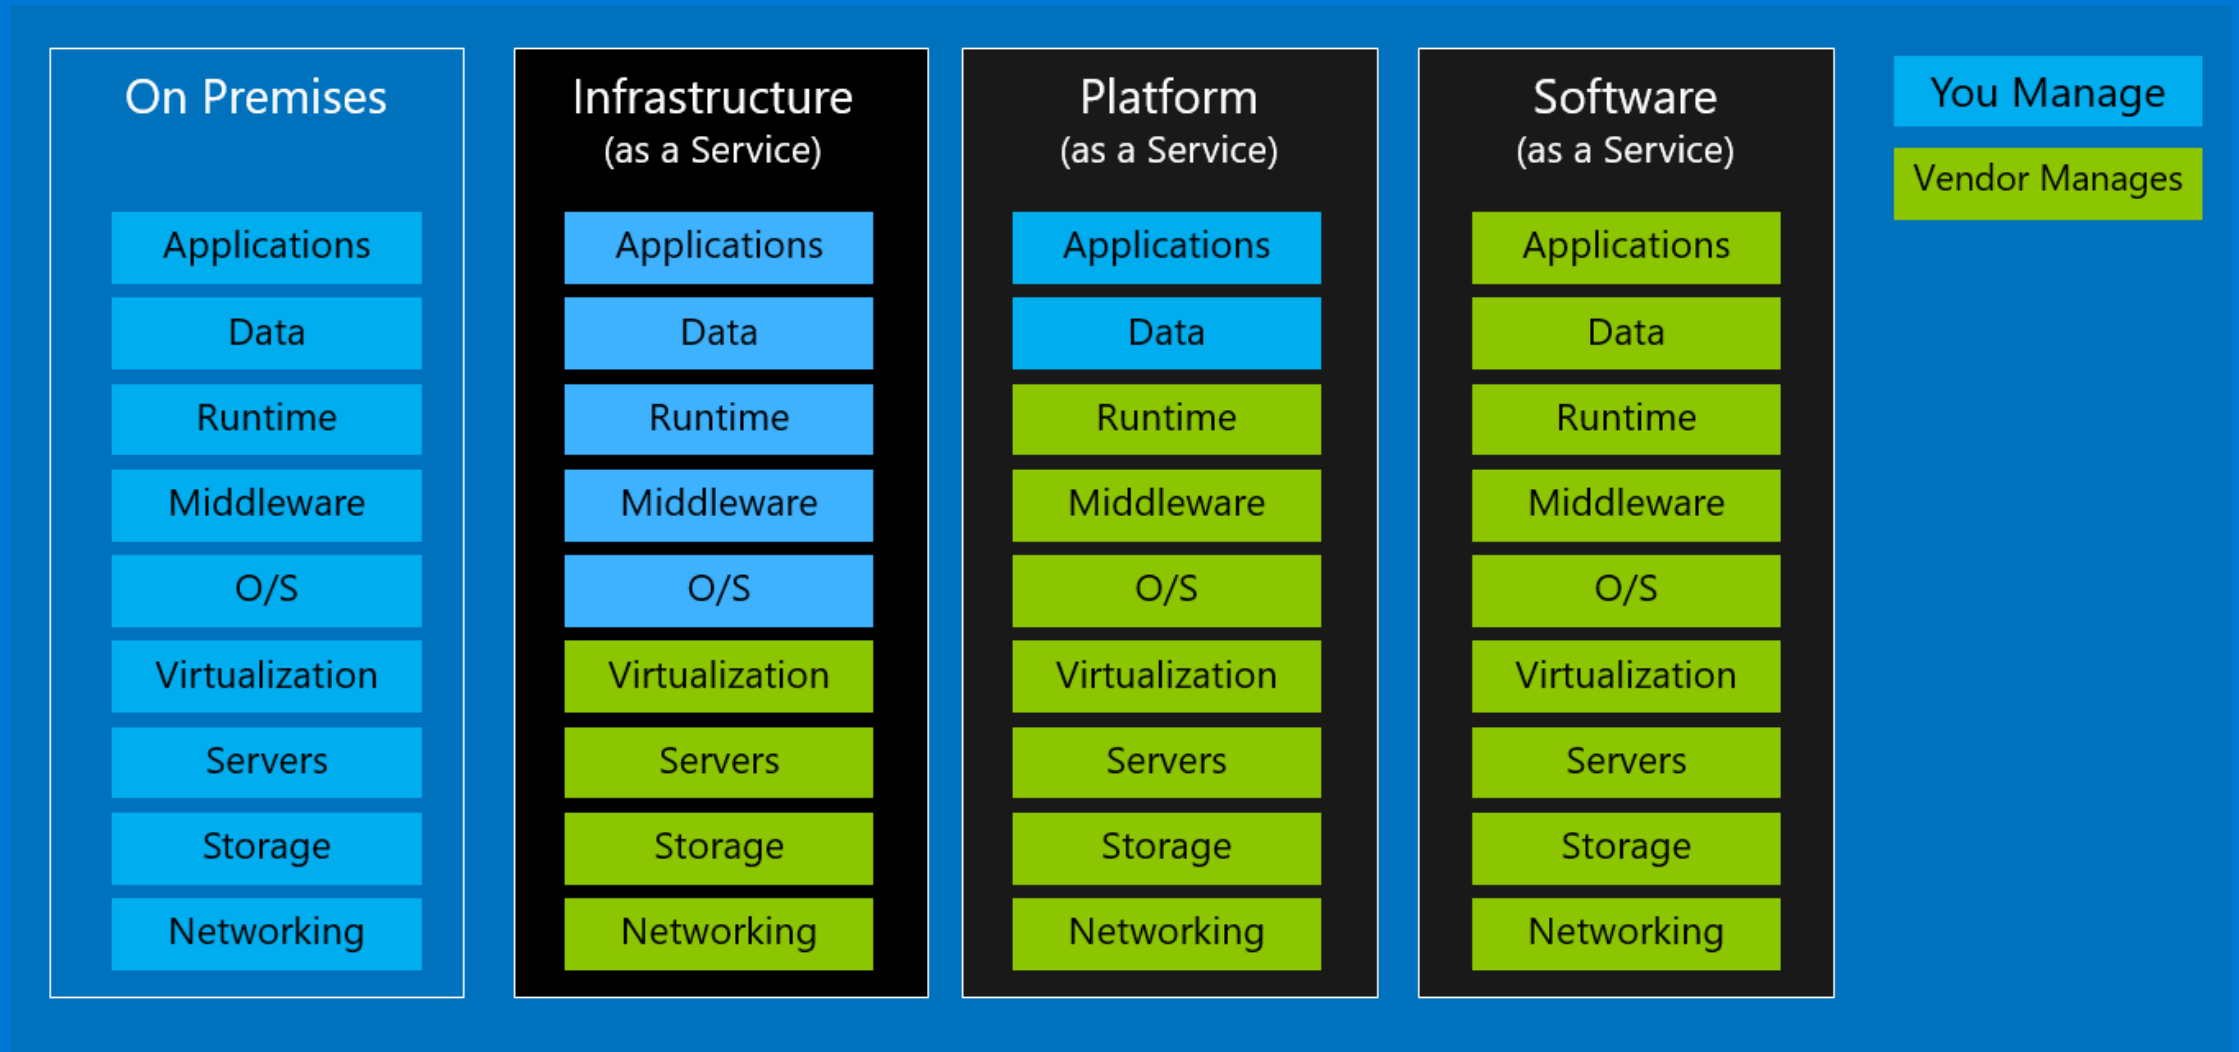
# Fluent Benchmark



# Fluent Benchmark




# Level Set: Cloud 101



# Application types in HPC



## Embarrassingly parallel:

- Applications do not communicate
- May share common store & data
- May have dependencies
- E.g. Monte Carlo simulations, transcoding, rendering



## Tightly coupled:

- Applications communicate; mainly use MPI
- Requires low latency, high bandwidth networking for scale
- E.g. car crash simulation, fluid dynamics, AI training






# Azure Big Compute – Solution Architectures


## Self-managed / Hybrid

## Fully-managed

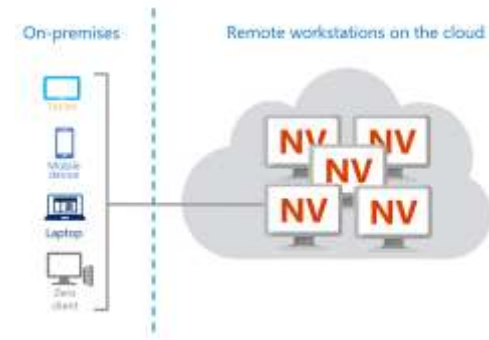
### Cloud burst




### Cluster on the cloud




### Remote Workstations



### HPC as a service

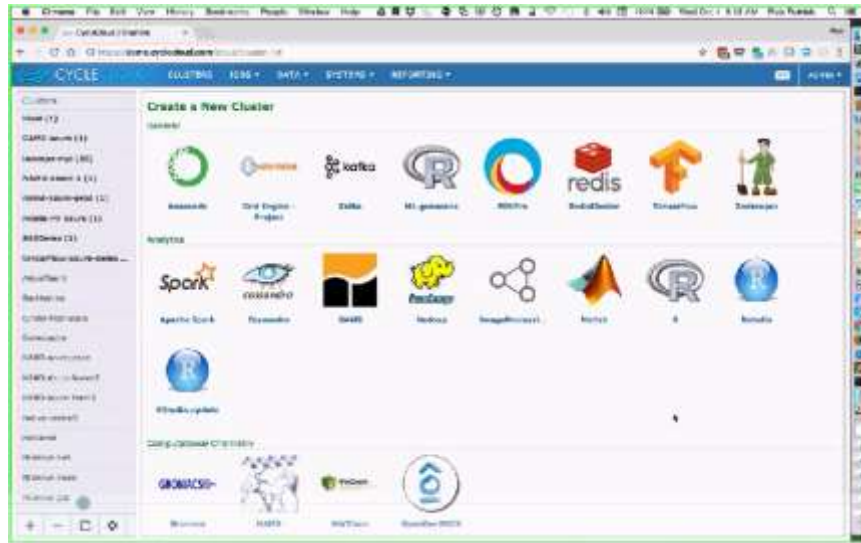


... and HPC is just **one building block** ... (150+ services in Azure, including:)

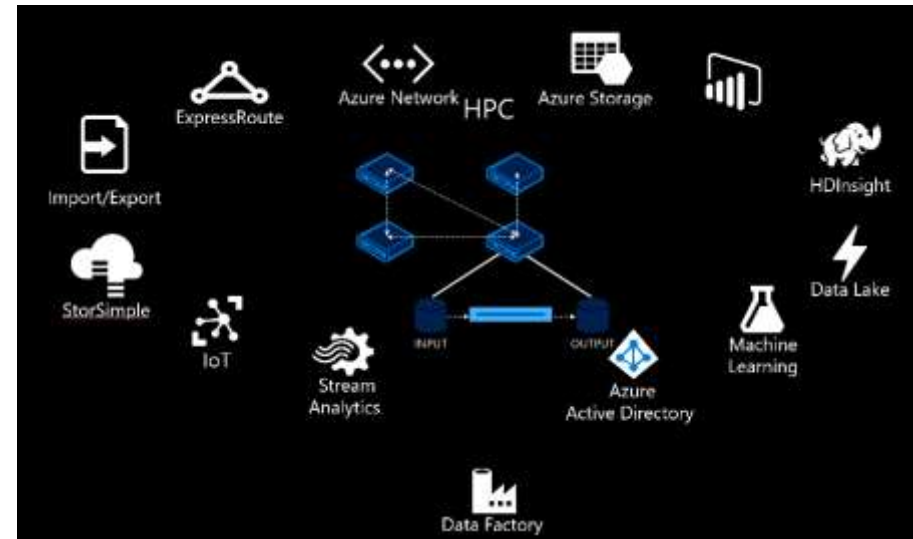


# Next Generation HPC – thinking differently with public cloud

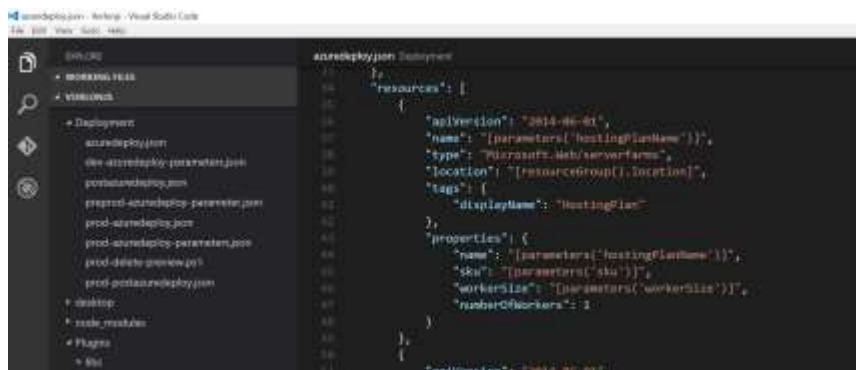
“A cluster (cloud!) for every HPC workload”



“HPC is one building block”



“HPC Infrastructure as code”



“AI, ML, Data Analytics”

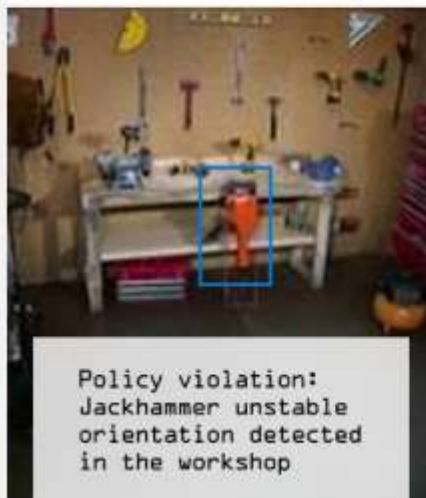


# Build in intelligence with AI: our Cognitive Services API

## Vision

From faces to feelings, allow apps to understand images and video

*Show me what is in the image*



## Speech

Hear and speak to users by filtering noise, identifying speakers, and understanding intent

*Convert this text to speech please...*



## Language

Process text and learn how to recognize what users want

*Play today's conference call...*

### Natural Language Processing

Intent: PlayCall  
Content: Customer#  
DateTime.date: today



Now Playing

11/29/2016 Customer Call

## Knowledge

Tap into rich knowledge amassed from the web, academia, or your own data

*Top publications in AI...*

### Genetic Algorithms in Search, Optimization Learning

1985, David E Goldberg

Cited 18,910 times

### Induction of Decision Trees

1986, Machine Learning

J R Quinlan

decision tree expert system machine learning compute

Cited 4,819 times

### Outline of a New Approach to the Analysis of Systems and Decision Processes

1973, IEEE Transactions on Systems, Man, and Cybernetics, Issn Lotfi A Zadeh (University Of California Berkeley)

fuzzy set complex systems artificial neural network computer machine learning artificial intelligence computer science

Cited 2,602 times

## Search

Locate relevant information among billions of web pages, images, videos, and news with Bing APIs

*Fraud prevention results...*

### How ID theft victims' in

nment documents or benefits

Credit card fraud

Phone or utilities fraud

Bank fraud 8.1%

Employment-related fraud 4.8%

Attempted identity theft 4.8%

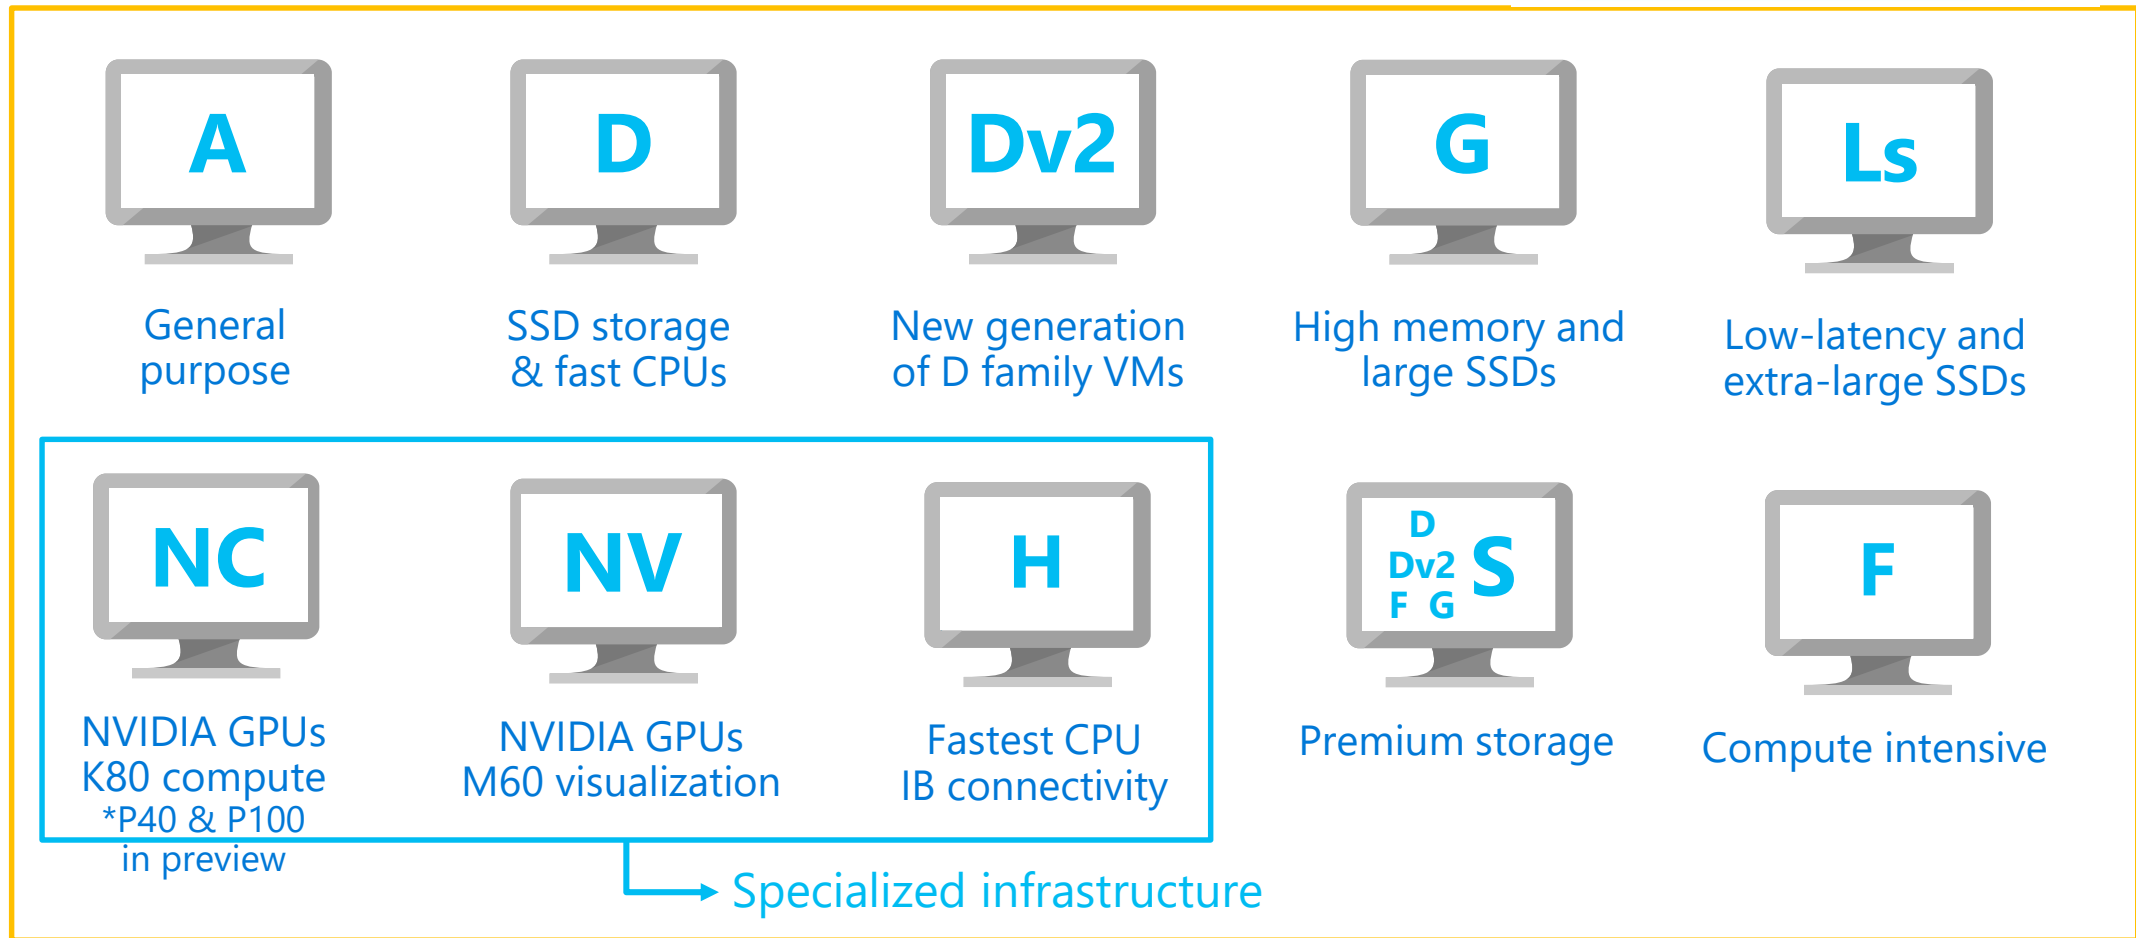
Loan fraud 4.4%

Other

sion Consumer Sentinel Network Data Book,

# Hardware capabilities

## VM sizes





# No-compromise HPC and AI VMs



- Up to 16 cores, 3.2 GHz E5-2667 V3 Haswell processor
- Up to 224 GiB DDR4 memory
- FDR InfiniBand (56 Gbps, 2.6 microsecond latency)
- 2 TB of local SSD



- Up to 4 NVIDIA Tesla K80 GPUs
- Up to 24 cores
- Up to 224 GiB memory
- Up to 1440 GiB of local SSD
- FDR InfiniBand



- Up to 4 NVIDIA Tesla M60 GPUs
- Up to 24 cores
- Up to 224 GiB memory
- Up to 1440 GiB of local SSD



- Up to 4 NVIDIA Pascal P40 GPUs
- Up to 24 cores
- Up to 448 GiB memory
- Up to 3 TB of local SSD
- FDR InfiniBand



- Up to 4 NVIDIA Pascal P100 GPUs
- Up to 24 cores
- Up to 448 GiB memory
- Up to 3 TB of local SSD
- FDR InfiniBand

## Clusters

blast (1)

GAMS-azure (1)

lammps-mpi (65)

NAMD-azure-1 (1)

namd-azure-prod (2)

nvidia-ml-azure (1)

tensorflow-azure-demo ...

AnsysFluent

Backtesting

condor-high-scale

DemoLustre

NAMD-azure-jason

NAMD-Azure-Jason-2

NAMD-azure-Team2

nat-us-central1

nvidia-ml

RCensus-aws

RCensus-azure

RCensus-gcp

sge-azure-benchmark



## lammps-mpi

Terminate

Edit

Share

Refresh

State **Started** at 6/20/17 11:19 AM (up 105d 19h 49m)

Nodes **1** ready, **64** acquiring

Scalesets **1** configuring

Size **1** instance, **16** cores (\$2.14 per hour)

Usage **806.8** core-hours (~\$52) in the last 24 hours

Alerts [Create new alert](#)

### Nodes Arrays Azure Scalesets

View: Template

Actions

Search

Template	Nodes	Cores	Status	Last Message
master	1	16		...
execute	64	1024	Allocating nodes	

View: Details

Show Detail

Edit

Connect

Actions

Search

Name	Type	Status	Cores	Zone	Host Name	UpTime
execute-1	Standard_H16mr	Acquiring	16	...	10.8.0.9	
execute-10	Standard_H16mr	Acquiring	16	...	10.8.0.10	
execute-11	Standard_H16mr	Acquiring	16	...	10.8.0.21	
execute-12	Standard_H16mr	Acquiring	16	...	10.8.0.33	
execute-13	Standard_H16mr	Acquiring	16	...	10.8.0.46	
execute-14	Standard_H16mr	Acquiring	16	...	10.8.0.58	
execute-15	Standard_H16mr	Acquiring	16	...	10.8.0.69	
execute-16	Standard_H16mr	Acquiring	16	...	10.8.0.74	
execute-17	Standard_H16mr	Acquiring	16	...	10.8.0.75	

Show: Active Instances by MachineType



Show Detail

Search

Time	Message
7:06 AM	Added 64 nodes of type execute to cluster lammps
9/28/17 9:20 AM	Node execute-6 in cluster lammps-mpi finished sta
9/28/17 9:19 AM	8 nodes in cluster lammps-mpi finished starting
9/28/17 9:19 AM	4 nodes in cluster lammps-mpi finished starting
9/28/17 9:19 AM	2 nodes in cluster lammps-mpi finished starting
9/28/17 9:19 AM	Node execute-11 in cluster lammps-mpi finished s
9/28/17 9:12 AM	Added 16 nodes of type execute to cluster lammps
9/19/17 5:44 AM	Node execute-25 in cluster lammps-mpi finished s
9/19/17 5:43 AM	3 nodes in cluster lammps-mpi finished starting
9/19/17 5:43 AM	2 nodes in cluster lammps-mpi finished starting
9/19/17 5:43 AM	2 nodes in cluster lammps-mpi finished starting
9/19/17 5:43 AM	7 nodes in cluster lammps-mpi finished starting
9/19/17 5:42 AM	4 nodes in cluster lammps-mpi finished starting
9/19/17 5:42 AM	5 nodes in cluster lammps-mpi finished starting
9/19/17 5:42 AM	2 nodes in cluster lammps-mpi finished starting
9/19/17 5:42 AM	4 nodes in cluster lammps-mpi finished starting
9/19/17 5:41 AM	4 nodes in cluster lammps-mpi finished starting
9/19/17 5:41 AM	6 nodes in cluster lammps-mpi finished starting
9/19/17 5:41 AM	4 nodes in cluster lammps-mpi finished starting
9/19/17 5:41 AM	3 nodes in cluster lammps-mpi finished starting
9/19/17 5:40 AM	6 nodes in cluster lammps-mpi finished starting
9/19/17 5:40 AM	3 nodes in cluster lammps-mpi finished starting

ChromeFileEditViewHistoryBookmarksPeopleWindowHelp

CycleCloud | Clusters

Rob

https://demo.cyclecloud.com/cloud/cluster\_list

CYCLECLOUD

CLUSTERSJOBS ▾DATA ▾SYSTEMS ▾REPORTING ▾

ADMIN ▾

Clusters

blast (1)

GAMS-azure (1)

lammps-mpi (65)

NAMD-azure-1 (1)

namd-azure-prod (1)

nvidia-ml-azure (1)

SGEDemo (1)

tensorflow-azure-demo ...

AnsysFluent

Backtesting

condor-high-scale

DemoLustre

NAMD-azure-jason

NAMD-Azure-Jason-2

NAMD-azure-Team2

nat-us-central1

nvidia-ml

RCensus-aws

RCensus-azure

RCensus-gcp

Create a New Cluster

General

Anaconda

Grid Engine - Project

Kafka

ML-genomics

PBSPro

RedisCluster

Tensorflow

Zookeeper

Analytics

Apache Spark

Cassandra

GAMS

Hadoop

ImageProcessi...

Matlab

R

Rstudio

RStudio-update

Computational Chemistry

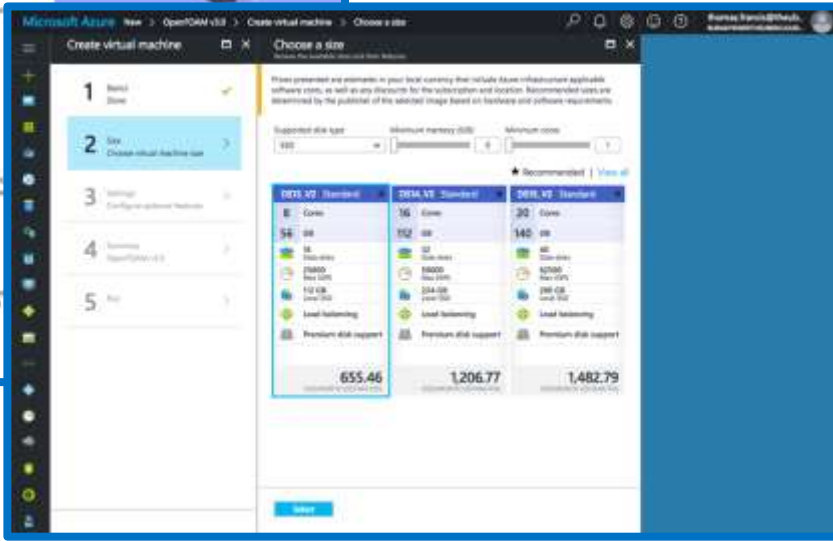
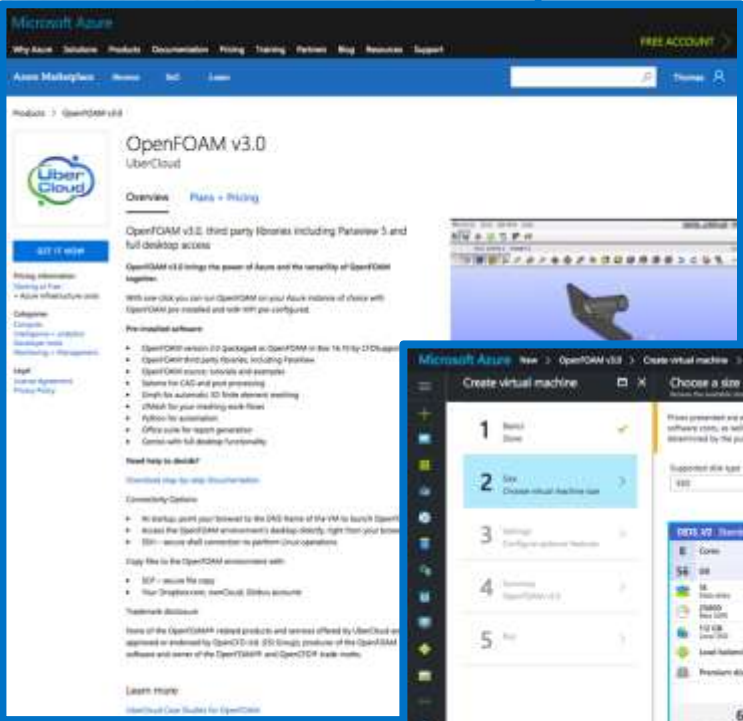
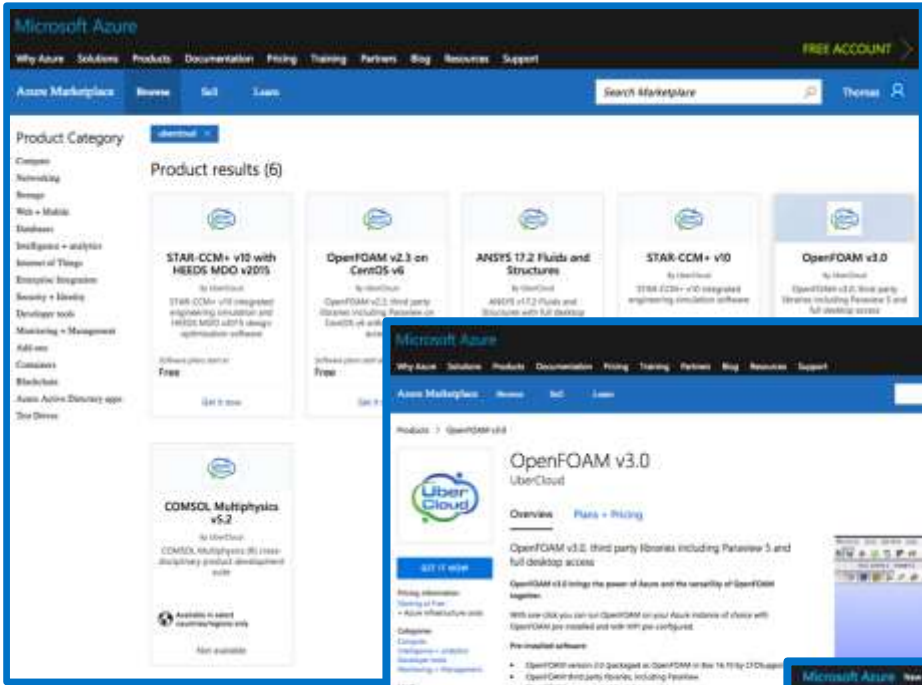
Gromacs

NAMD

NWChem

OpenEye ROCS

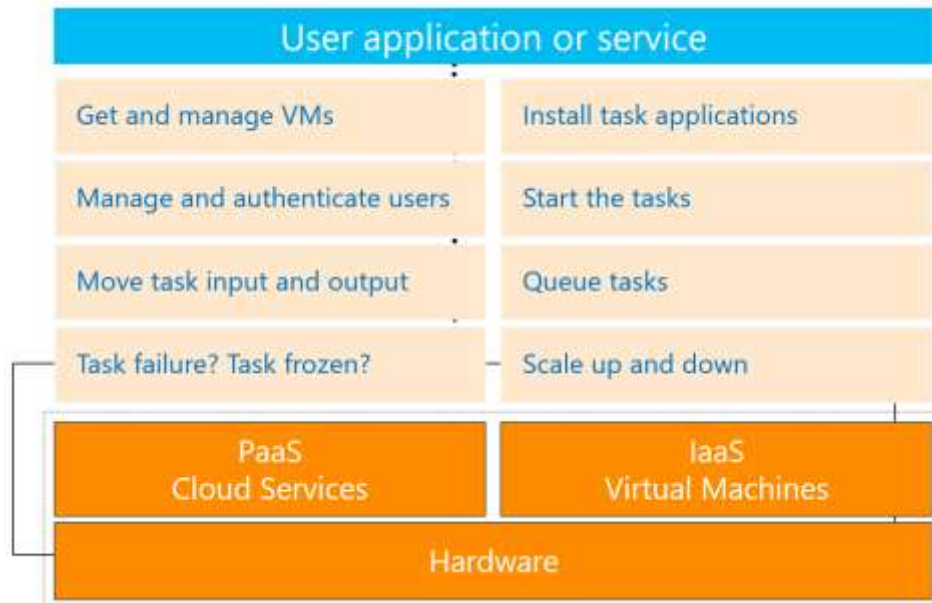




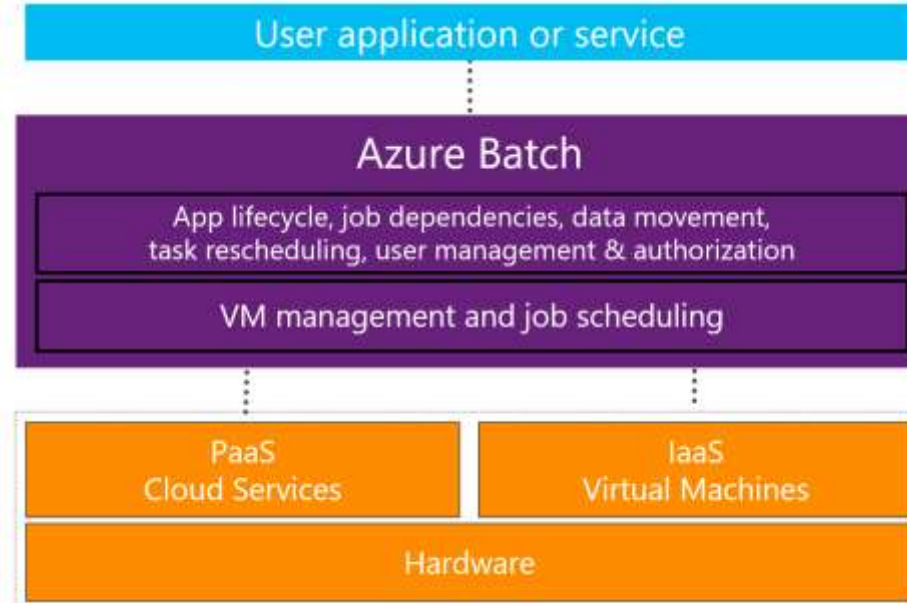


# Azure Batch – Concepts

Before....



After...





- Code directly against this service (it's an API, and it's free)
- Let Azure Batch manage the resources & scheduling
- Use any resources (any VM instance, GPUs, RDMA/MPI, etc)
- Popular with developers, software vendors (ISVs)
- Create "**Pools**", submit "**Jobs**" made up of "**Tasks**"
- Mix & match, dynamically resize, include Low Pri VMs to save!



# Elasticity & Scale

What would you do with 100,000 cores? -  
Big compute at global scale

<https://azure.microsoft.com/en-us/blog/what-would-you-do-with-100000-cores-big-compute-at-global-scale/>




<https://blogs.endjin.com/2015/07/spinning-up-16000-a1-virtual-machines-on-azure-batch/>

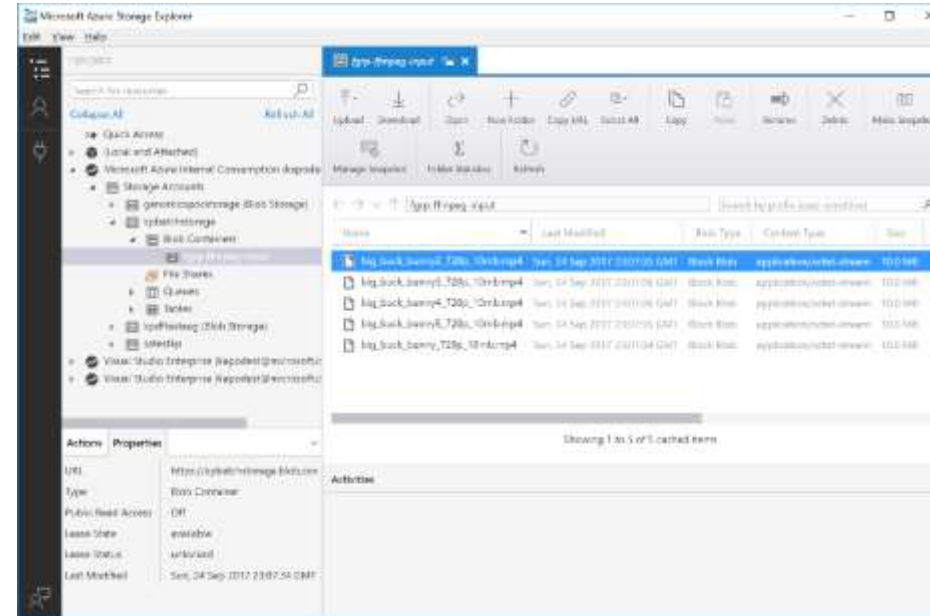


# Azure Batch – Lab setup

## Azure Batch Labs



## Azure Storage Explorer



Batch AI Lab:

<https://github.com/azurebigcompute/CloudWorkshops/blob/master/BatchAIWorkshop.md>

Batch with FFMPEG Lab:

<https://github.com/azurebigcompute/Labs/tree/master/Azure%20Batch%20Masterclass%20Labs>

# Some Links to get started!

Azure Big Compute (Labs, Solutions, Templates) - <https://github.com/azurebigcompute>

Azure Pricing Calculator - <https://azure.microsoft.com/en-us/pricing/calculator/>

Azure Documentation - <https://docs.microsoft.com/en-us/azure/>

Code samples - <https://azure.microsoft.com/en-us/resources/samples/>

Templates - <https://azure.microsoft.com/en-us/resources/templates/>